ML: RANDOM FORESTS AND GRADIENT BOOSTING

MGMT 675
AI-Assisted Financial Analysis
Kerry Back

RICE | BUSINESS
Jones Graduate School of Business

MORE MODELS

- Decision tree

- Random forest

- Boosted trees

- Neural networks

DECISION TREE

- Split dataset successively into subsets. Within each subset, $\hat{y} = $ mean of subset. Calculate MSE.

- Split on a single variables being above or below a threshold.

- Choose variable and threshold so that MSE will be as small as possible after the split.

- After each split, make further splits of all of the new subsets into even smaller subsets, for a specified number of times (# splits = depth).

- The prediction for any observation is the mean target value in its final group (leaf).

EXAMPLE

- Ask Julius to read ml1.xlsx.

- Ask Julius to fit a decision tree regressor with y1 as the target using all of the data as training data. Ask Julius to plot the tree.

RANDOM FOREST

- Generate random datasets of the same size as the original.

- Create the random datasets by randomly drawing rows from the original with replacement.

- Fit a decision tree to each random dataset.

- The prediction for any observation is the average of the predictions of the various trees.

- Randomization helps to avoid overfitting.
- Also control overfitting through:
  - max_depth = maximum number of times to split in each tree
  - max_features = number of features to look at when deciding how to split (a subset of features of that size is randomly chosen for each split)

GRADIENT BOOSTING

- Fit a decision tree.

- Look at its errors. Fit a new decision tree to predict the errors.

- New prediction is original plus a fraction of the prediction of original's error (fraction = learning rate).

- Look at the errors of the new predictions. Fit a new decision to predict these errors.

- Continue …

EXAMPLES

- Ask Julius to train and test a random forest regressor to predict y1 in ml1.xlsx.

- Ask Julius to use GridSearchCV to find the best max_depth in (5, 10, 15, 20).

- Ask Julius to train and test a gradient boosting regressor to predict y1 in ml1.xlsx.

- Ask Julius to use GridSearchCV to find the best learning rate in (0.001, 0.005, 0.01, 0.05, 0.1).

INTERPRETING MODELS: SHAPLEY VALUES

- The Shapley value for a feature at an observation is a measure of how much that feature contributed to the prediction at that observation.

- A summary of Shapley values is a bar chart showing the mean absolute contribution of each feature (mean across observations).

- A Shapley scatter plot for a feature plots all of the observations with the feature's value on the x axis and the feature's contribution to the prediction on the y axis.

- Ask Julius to create a summary plot of the Shapley values for the random forest regressor with the best max_depth.

- Ask Julius to create a scatter plot of the Shapley values for the x1 feature.

- Ask Julius to create a scatter plot of the Shapley values for another feature.

NEURAL NETWORKS

- The most common type of neural network is a multi-layer perceptron.

-