

ML - RANDOM FORESTS AND GRADIENT BOOSTING

MGMT 675

AI-Assisted Financial Analysis

Kerry Back



RICE | BUSINESS
Jones Graduate School of Business

OUTLINE

- Decision tree
- Random forest and gradient boosting
- Shapley values
- House price application

DECISION TREE

- Split dataset successively into subsets. Within each subset, \hat{y} = mean of subset. Calculate MSE.
- Split on a single variable being above or below a threshold.
- Choose variable and threshold so that MSE will be as small as possible after the split.

- After each split, make further splits of all of the new subsets into even smaller subsets, for a specified number of times (# splits = depth).
- The prediction for any observation is the mean target value in its final group (leaf).

EXAMPLE

- Ask Julius to read ml1.xlsx.
- Ask Julius to fit a decision tree regressor with y1 as the target using all of the data as training data. Ask Julius to plot the tree.

RANDOM FOREST AND GRADIENT BOOSTING

RANDOM FOREST

- Generate random datasets of the same size as the original.
- Create the random datasets by randomly drawing rows from the original with replacement.
- Fit a decision tree to each random dataset.
- The prediction for any observation is the average of the predictions of the various trees.

- Randomization helps to avoid overfitting.
- Also control overfitting through:
 - `max_depth` = maximum number of times to split in each tree
 - `max_features` = number of features to look at when deciding how to split (a subset of features of that size is randomly chosen for each split)

GRADIENT BOOSTING

- Fit a decision tree.
- Look at its errors. Fit a new decision tree to predict the errors.
- New prediction is original plus a fraction of the prediction of original's error (fraction = learning rate).
- Look at the errors of the new predictions. Fit a new decision to predict these errors.
- Continue ...

EXAMPLES

- Ask Julius to train and test a random forest regressor to predict `y1` in `ml1.xlsx`.
- Ask Julius to use `GridSearchCV` to find the best `max_depth` in `(5, 10, 15, 20)`.
- Ask Julius to train and test a gradient boosting regressor to predict `y1` in `ml1.xlsx`.
- Ask Julius to use `GridSearchCV` to find the best learning rate in `(0.001, 0.005, 0.01, 0.05, 0.1)`.

INTERPRETING MODELS: SHAPLEY VALUES

- The Shapley value for a feature at an observation is a measure of how much that feature contributed to the prediction at that observation.
- A summary of Shapley values is a bar chart showing the mean absolute contribution of each feature (mean across observations).
- A Shapley scatter plot for a feature plots all of the observations with the feature's value on the x axis and the feature's contribution to the prediction on the y axis.

- Ask Julius to create a summary plot of the Shapley values for the random forest regressor with the best max_depth.
- Ask Julius to create a scatter plot of the Shapley values for the x1 feature.
- Ask Julius to create a scatter plot of the Shapley values for another feature.

HOUSE PRICE APPLICATION (TBD)